


Slides for Chapter 15: Coordination and Agreement



From **Coulouris, Dollimore, Kindberg and Blair**

**Distributed Systems:
Concepts and Design**

Edition 5, © Addison-Wesley 2012

Overview of Chapter


- Introduction
- Distributed mutual exclusion
- Elections
- Coordination and agreement in group communication (skip)
- Consensus and related problems (skip)

Introduction

Covers two areas:

- Coordinating actions in a distributed system
- Distributed processes agreeing on a result value
- Assumes reliable communication channels for simplicity (failure is masked by a reliable communication protocol)
- Detecting that a process has failed can be reliable or unreliable
- Use timeouts
- Unreliable: replies *unsuspected* or *suspected*
- Reliable: replies *unsuspected* or *failed*

Figure 15.1
A network partition



Overview of Chapter

- Introduction
- Distributed mutual exclusion
- Elections
- Coordination and agreement in group communication
- Consensus and related problems

Distributed mutual exclusion

Known as *critical section* problem:

- Only one process can be in critical section – the process with the *token* that allows access to the resource
- Operations include the following:
- `enter()` – requests access; can be granted or blocked
- `resourceAccess()` – access resource in critical section
- `exit()` – leave critical section – other processes may now enter

Conditions:

- ME1 (safety) – at most one process in critical section
- ME2 (liveness) – requests eventually succeed
- ME3 (ordering) – requests follow happened-before relationship

Distributed mutual exclusion

Various algorithms:

- Central server algorithm
- Ring-based algorithm
- Algorithm using multicast and logical clocks
- Voting algorithm
- Others

Figure 15.2
Server managing a mutual exclusion token for a set of processes




Figure 15.3
A ring of processes transferring a mutual exclusion token




Figure 15.4

Ricart and Agrawala's algorithm

On initialization

$state := \text{RELEASED};$

To enter the section

$state := \text{WANTED};$

Multicast *request* to all processes;

$T := \text{request's timestamp};$

Wait *until* (number of replies received = $(N - 1)$);

$state := \text{HELD};$

} request processing deferred here

On receipt of a request $\langle T_i, p_i \rangle$ at p_j ($i \neq j$)

if ($state = \text{HELD}$ or ($state = \text{WANTED}$ and $(T, p_j) < (T_i, p_i)$))

then

queue *request* from p_i without replying;

else

reply immediately to p_i ;

end if

To exit the critical section

$state := \text{RELEASED};$

reply to any queued requests;

Figure 15.5
Multicast synchronization




Figure 15.6

Maekawa's algorithm – part 1

On initialization

state := RELEASED;

voted := FALSE;

For p_i *to enter the critical section*

state := WANTED;

Multicast *request* to all processes in V_i ;

Wait until (number of replies received = K);

state := HELD;

On receipt of a request from p_i *at* p_j

if (*state* = HELD or *voted* = TRUE)

then

 queue *request* from p_i without replying;

else

 send *reply* to p_i ;

voted := TRUE;

end if

For p_i *to exit the critical section*

state := RELEASED;

Multicast *release* to all processes in V_i ;

On receipt of a release from p_i *at* p_j

if (queue of requests is non-empty)

then

 remove head of queue – from p_k , say;

 send *reply* to p_k ;

voted := TRUE;

else

voted := FALSE;

end if

Overview of Chapter

- Introduction
- Distributed mutual exclusion
- Elections
- Coordination and agreement in group communication
- Consensus and related problems

Elections

Election algorithms:

- Used to choose a particular process for a role – for example, to choose a central server for distributed algorithms that require a central server
- Process can call an election; for example, if it detects that central server has failed

Requirements:


- E1: (safety) only one process is elected – each participant sets elected_i either to P or to undefined if it does not know the elected process yet (P will be the participant with the largest process id)
- E2: (liveness) all processes participate and either set elected_i to the elected process or crash

Elections

Election algorithms:

- Ring-based algorithm
- Bully algorithm

Figure 15.7
A ring-based election in progress




Note: The election was started by process 17.
The highest process identifier encountered so far is 24.
Participant processes are shown in a darker colour

Figure 15.8

The bully algorithm

The election of coordinator p_2 ,
after the failure of p_4 and then p_3



Overview of Chapter

- Introduction
- Distributed mutual exclusion
- Elections
- Coordination and agreement in group communication (skip)
- Consensus and related problems

Figure 15.9

Reliable multicast algorithm

On initialization

Received := {};

For process p to R-multicast message m to group g

B-multicast(g, m); // $p \in g$ is included as a destination

On B-deliver(m) at process q with $g = \text{group}(m)$

if ($m \notin \text{Received}$)

then

Received := *Received* \cup {*m*};

if ($q \neq p$) then B-multicast(g, m); end if

R-deliver m;

end if

Figure 15.10

The hold-back queue for arriving multicast messages

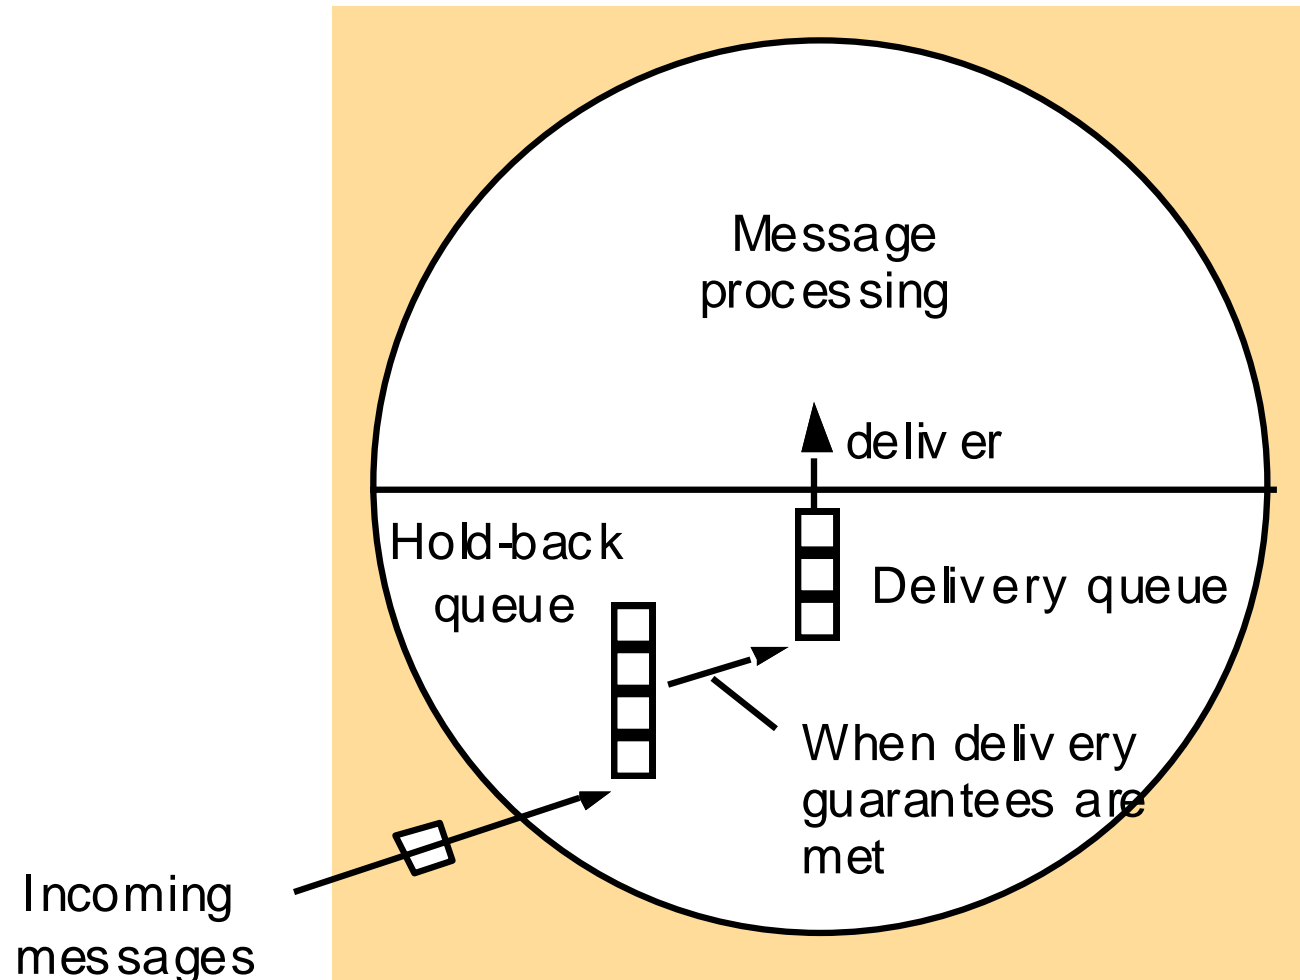


Figure 15.11

Total, FIFO and causal ordering of multicast messages

Notice the consistent ordering of totally ordered messages T_1 and T_2 , the FIFO-related messages F_1 and F_2 and the causally related messages C_1 and C_3 – and the otherwise arbitrary delivery ordering of messages.




Figure 15.12

Display from bulletin board program

Bulletin board: os.interesting		
Item	From	Subject
23	A.Hanlon	Mach
24	G.Joseph	Microkernels
25	A.Hanlon	Re: Microkernels
26	T.L'Heureux	RPC performance
27	M.Walker	Re: Mach
end		

Figure 15.13

Total ordering using a sequencer

1. Algorithm for group member p

On initialization: $r_g := 0$;

To TO-multicast message m to group g

B-multicast($g \cup \{\text{sequencer}(g)\}$, $\langle m, i \rangle$);

On B-deliver($\langle m, i \rangle$) with $g = \text{group}(m)$

Place $\langle m, i \rangle$ in hold-back queue;

On B-deliver($m_{\text{order}} = \langle \text{"order"}, i, S \rangle$) with $g = \text{group}(m_{\text{order}})$

wait until $\langle m, i \rangle$ in hold-back queue and $S = r_g$;

TO-deliver m ; // (after deleting it from the hold-back queue)

$r_g = S + 1$;

2. Algorithm for sequencer of g

On initialization: $s_g := 0$;

On B-deliver($\langle m, i \rangle$) with $g = \text{group}(m)$

B-multicast(g , $\langle \text{"order"}, i, s_g \rangle$);

$s_g := s_g + 1$;

Figure 15.14
The ISIS algorithm for total ordering

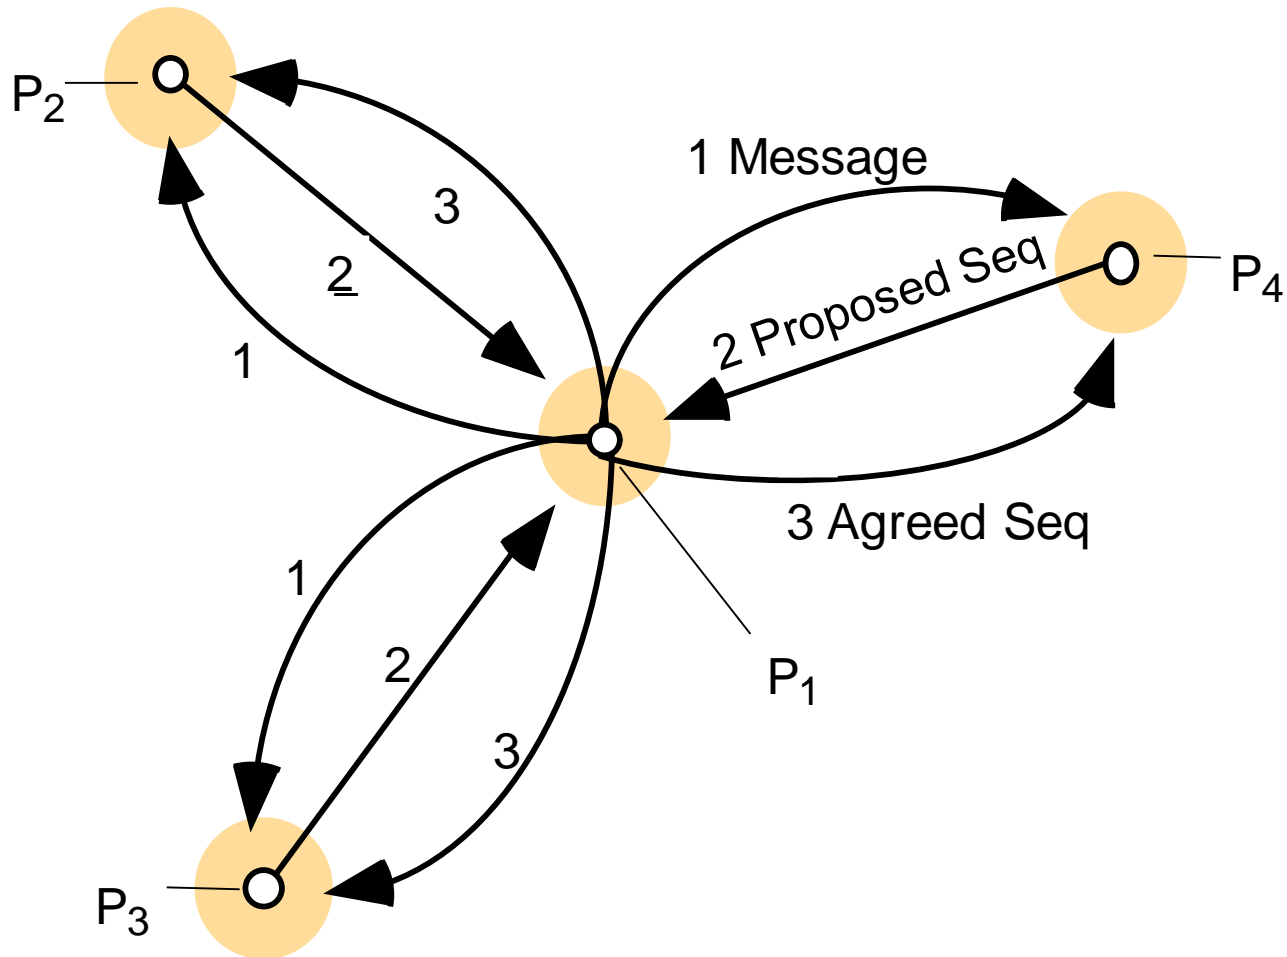


Figure 15.15

Causal ordering using vector timestamps

Algorithm for group member p_i ($i = 1, 2, \dots, N$)

On initialization

$V_i^g[j] := 0$ ($j = 1, 2, \dots, N$);

To CO-multicast message m to group g

$V_i^g[i] := V_i^g[i] + 1$;

$B\text{-multicast}(g, \langle V_i^g, m \rangle)$;

On $B\text{-deliver}(\langle V_j^g, m \rangle)$ from p_j , with $g = \text{group}(m)$

place $\langle V_j^g, m \rangle$ in hold-back queue;

wait until $V_j^g[j] = V_i^g[j] + 1$ and $V_j^g[k] \leq V_i^g[k]$ ($k \neq j$);

$CO\text{-deliver } m$; // after removing it from the hold-back queue

$V_i^g[j] := V_i^g[j] + 1$;

Figure 15.16
Consensus for three processes

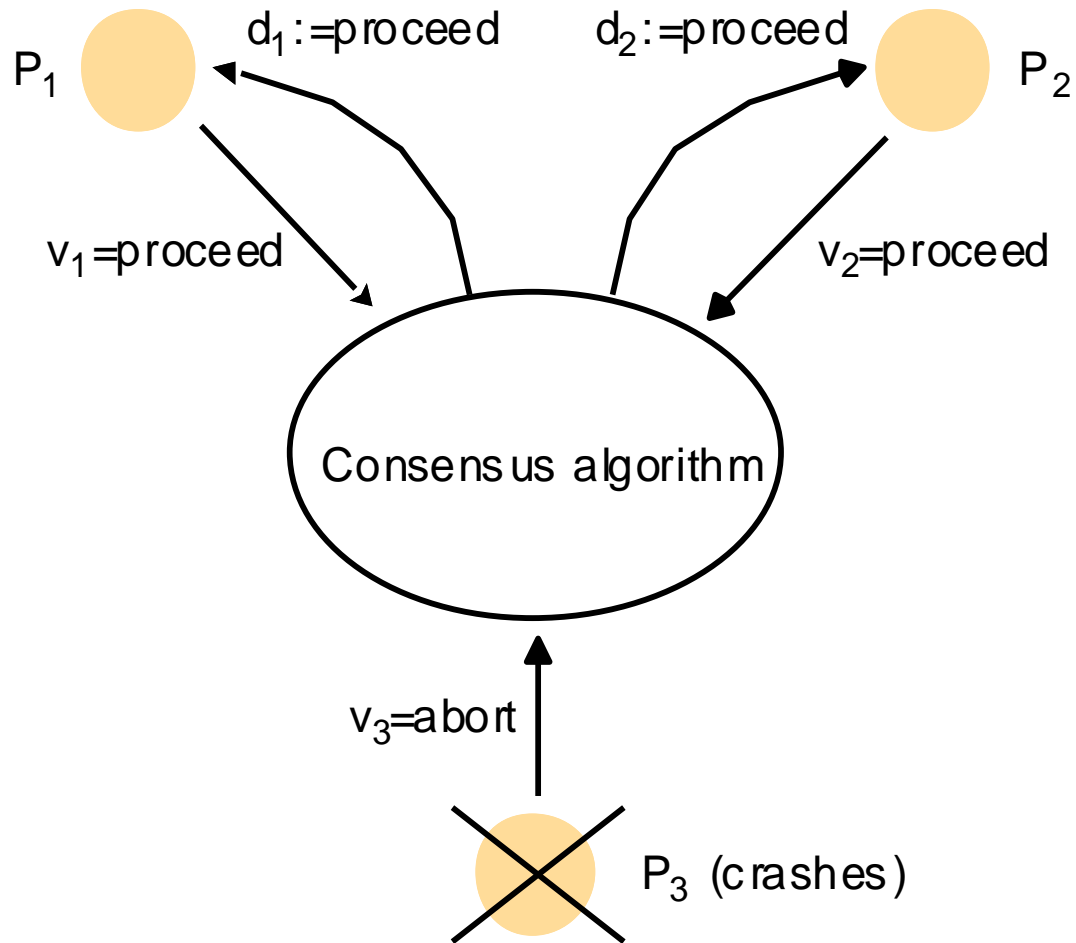


Figure 15.17

Consensus in a synchronous system

Algorithm for process $p_i \in g$; algorithm proceeds in $f + 1$ rounds

On initialization

$Values_i^1 := \{v_i\}; Values_i^0 = \{\};$

In round r ($1 \leq r \leq f + 1$)

$B\text{-multicast}(g, Values_i^r - Values_i^{r-1});$ // Send only values that have not been sent

$Values_i^{r+1} := Values_i^r;$

while (in round r)

{

On B-deliver(V_j) from some p_j



$Values_i^{r+1} := Values_i^{r+1} \cup V_j;$

}

After $(f + 1)$ rounds



Assign $d_i = \text{minimum}(Values_i^{f+1});$

Figure 15.18
Three Byzantine generals



Faulty processes are shown coloured

Figure 15.19
Four Byzantine generals



Faulty processes are shown coloured